

# Project #1

My solution

# Projection-based method

- Items are listed in lexicographic order
- Let  $P$  and  $DB(P)$  be a node's pattern and its associated projected database.
- Mining is performed by recursively calling this function:
  - $TP(P, DB(P))$ 
    1. Determine the frequent items in  $DB(P)$ , and denote them by  $E(P)$ .
    2. Eliminate from  $DB(P)$  any items not in  $E(P)$ .
    3. For each item  $x$  in  $E(P)$ , call  $TP(Px, DB(Px))$ .

# What is the right data structure for storing the projected database?


- Operations that needs to support:
  - Determining the frequency of each item.
  - Projection:
    - Extract the transactions and their subsets of items that support a given pattern.
- What should be the upper bound on the complexity of each of these operations?

Dense matrices are stored in the computer memory by using two-dimensional arrays. For example, a matrix with  $n$  rows and  $m$  columns, is stored using a  $n \times m$  array of real numbers. However, using the same two-dimensional arrays to store sparse matrices has two very important drawbacks. First, since most of the entries in the sparse matrix are zero, this storage scheme wastes a lot of memory. Second, computations involving sparse matrices often need to operate only on the non-zero entries of the matrix. Use of dense storage format makes it harder to locate these non-zero entries. For these reasons sparse matrices are stored using different data structures.

The compressed storage format (CSR) is a widely used scheme for storing sparse matrices. In the CSR format, an  $n \times n$  sparse matrix  $A$  that has  $k$  non-zero entries is stored using three

arrays: two integer arrays **rowptr** and **colind**, and one array of real values **values**. The array **rowptr** is of size  $n + 1$ , and the other two arrays are each of size  $k$ . The array **colind** stores the column-indices of the non-zero entries in  $A$ , and the array **values** stores the corresponding non-zero entries. In particular, the array **colind** stores the column-indices of the first row followed by the column-indices of the second row followed by the column-indices of the third row, and so on. Similarly, the array **values** stores the corresponding non-zero entries of the first row followed by the corresponding non-zero entries of the second row, and so on. The array **rowptr** is used to determine where the storage of a row starts and ends in the arrays **colind** and **values**.

In particular, the column-indices of row  $i$  are stored starting at **colind**[**rowptr**[ $i$ ]] and ending at (but not including) **colind**[**rowptr**[ $i+1$ ]]. Similarly, the values of the non-zero entries of row  $i$  are stored starting at **values**[**rowptr**[ $i$ ]] and ending at (but not including) **values**[**rowptr**[ $i+1$ ]]. Also note that the number of non-zero entries for row  $i$  is simply **rowptr**[ $i+1$ ]-**rowptr**[ $i$ ].



CSR format of a sample matrix

**Figure 0.11:** The compressed storage format for sparse matrices.

My naming of the CSR/CSC fields is different from the common way of doing things.

I use “row” as the prefix for the row-wise representation (rowptr, rowind, rowval).

I use “col” as the prefix for the column-wise representation (colptr, colind, colval).

```
/*-----  
 * The following data structure stores a sparse CSR format  
 *-----*/  
typedef struct gk_csr_t {  
    int32_t nrows, ncols;  
    → ssize_t *rowptr, *colptr;  
    → int32_t *rowind, *colind;  
    int32_t *rowids, *colids;  
    int32_t *rlabels, *clabels;  
    int32_t *rmap, *cmap;  
    → float *rowval, *colval;  
    float *rnorms, *cnorms;  
    float *rsums, *csums;  
    float *rsizes, *csizes;  
    float *rvols, *cvols;  
    float *rwgts, *cwgts;  
} gk_csr_t;
```

```

/*-----*/
/*! Data structures for use within this module */
/*-----*/
typedef struct {
    int minfreq; /* the minimum frequency of a pattern */
    int maxfreq; /* the maximum frequency of a pattern */
    int minlen; /* the minimum length of the requested pattern */
    int maxlen; /* the maximum length of the requested pattern */
    int titems; /* the initial range of the item space */

    /* the call-back function */
    void (*callback)(void *stateptr, int nitems, int *itemids, int ntrans, int *transids);
    void *stateptr; /* the user-supplied pointer to pass to the callback */

    /* workspace variables */
    int *rmarker;
    gk_ikv_t *cand;
} isparams_t;

```

The memory for  
those is allocated  
once and reused.



This is to keep track of the original item-IDs, since their numbering will change due to reordering.

```

/*****
/* The entry point of the frequent itemset discovery code */
*****/
void gk_find_frequent_itemsets(int ntrans, ssize_t *tranptr, int *tranind,
    int minfreq, int maxfreq, int minlen, int maxlen,
    void (*process_itemset)(void *stateptr, int nitems, int *itemids,
        int ntrans, int *transids),
    void *stateptr)
{
    ssize_t i;
    gk_csr_t *mat, *pmat;
    isparams_t params;
    int *pattern;

    /* Create the matrix */
    mat = gk_csr_Create();
    mat->nrows = ntrans;
    mat->ncols = tranind[gk_iargmax(tranptr[ntrans], tranind, 1)]+1;
    mat->rowptr = gk_zcopy(ntrans+1, tranptr,
        gk_zmalloc(ntrans+1, "gk_find_frequent_itemsets: mat.rowptr"));
    mat->rowind = gk_icopy(tranptr[ntrans], tranind,
        gk_imalloc(tranptr[ntrans], "gk_find_frequent_itemsets: mat.rowind"));
    mat->colids = gk_iincset(mat->ncols, 0,
        gk_imalloc(mat->ncols, "gk_find_frequent_itemsets: mat.colids"));

    /* Setup the parameters */
    params.minfreq = minfreq;
    params.maxfreq = (maxfreq == -1 ? mat->nrows : maxfreq);
    params.minlen = minlen;
    params.maxlen = (maxlen == -1 ? mat->ncols : maxlen);
    params.tnitems = mat->ncols;
    params.callback = process_itemset;
    params.stateptr = stateptr;
    params.rmarker = gk_ismalloc(mat->nrows, 0, "gk_find_frequent_itemsets: rmarker");
    params.cand = gk_ikvmalloc(mat->ncols, "gk_find_frequent_itemsets: cand");

    /* Perform the initial projection */
    gk_csr_CreateIndex(mat, GK_CSR_COL);
    pmat = itemsets_project_matrix(&params, mat, -1);
    gk_csr_Free(&mat);

    pattern = gk_imalloc(pmat->ncols, "gk_find_frequent_itemsets: pattern");
    itemsets_find_frequent_itemsets(&params, pmat, 0, pattern);

    gk_csr_Free(&pmat);
    gk_free((void **)&pattern, &params.rmarker, &params.cand, LTERM);
}

```

```

/*****
/*! The recursive routine for DFS-based frequent pattern discovery */
*****/
void itemsets_find_frequent_itemsets(isparams_t *params, gk_csr_t *mat,
    int preflen, int *prefix)
{
    ssize_t i;
    gk_csr_t *cmat;

    /* Project each frequent column */
    for (i=0; i<mat->ncols; i++) {
        prefix[preflen] = mat->colids[i];

        if (preflen+1 >= params->minlen)
            (*params->callback)(params->stateptr, preflen+1, prefix,
                mat->colptr[i+1]-mat->colptr[i], mat->colind+mat->colptr[i]);

        if (preflen+1 < params->maxlen) {
            cmat = itemsets_project_matrix(params, mat, i);
            itemsets_find_frequent_itemsets(params, cmat, preflen+1, prefix);
            gk_csr_Free(&cmat);
        }
    }
}

```



Why?

```
/******  
/*! This function projects a matrix w.r.t. to a particular column.  
It performs the following steps:  
- Determines the length of each column that is remaining.  
- Sorts the columns in increasing length.  
- Creates a column-based version of the matrix with the proper  
  column ordering.  
*/  
/******  
gk_csr_t *itemsets_project_matrix(isparams_t *params, gk_csr_t *mat, int cid)
```

```

gk_csr_t *itemsets_project_matrix(isparams_t *params, gk_csr_t *mat, int cid)
{
    ssize_t i, j, k, ii, pnnz;
    int nrows, ncols, pnrows, pncols;
    ssize_t *colptr, *pcolptr;
    int *colind, *colids, *pcolind, *pcolids, *rmarker;
    gk_csr_t *pmat;
    gk_ikv_t *cand;

    nrows = mat->nrows;
    ncols = mat->ncols;
    colptr = mat->colptr;
    colind = mat->colind;
    colids = mat->colids;

    rmarker = params->rmarker;
    cand = params->cand;

    /* Allocate space for the projected matrix based on what you know thus far */
    pmat = gk_csr_Create();
    pmat->nrows = pnrows = (cid == -1 ? nrows : colptr[cid+1]-colptr[cid]);

    /* Mark the rows that will be kept and determine the prowids */
    if (cid == -1) { /* Initial projection */
        gk_iset(nrows, 1, rmarker);
    }
    else { /* The other projections */
        for (i=colptr[cid]; i<colptr[cid+1]; i++)
            rmarker[colind[i]] = 1;
    }

    /* Determine the length of each column that will be left in the projected matrix */
    for (pncols=0, pnnz=0, i=cid+1; i<ncols; i++) {
        for (k=0, j=colptr[i]; j<colptr[i+1]; j++) {
            k += rmarker[colind[j]];
        }
        if (k >= params->minfreq && k <= params->maxfreq) {
            cand[pncols].val = i;
            cand[pncols++].key = k;
            pnnz += k;
        }
    }

    /* Sort the columns in increasing order */
    gk_ikvsorti(pncols, cand);
}

```

**pncols**: is the number of columns in the projected database.

**pnnz**: is the total number of items in the transactions of the projected database.

This is either 0 or 1.

```

/* Sort the columns in increasing order */
gk_ikvsorti(pncols, cand);

/* Allocate space for the remaining fields of the projected matrix */
pmat->ncols = pncols;
pmat->colids = pcolids = gk_imalloc(pncols, "itemsets_project_matrix: pcolids");
pmat->colptr = pcolptr = gk_zmalloc(pncols+1, "itemsets_project_matrix: pcolptr");
pmat->colind = pcolind = gk_imalloc(pnnz, "itemsets_project_matrix: pcolind");

/* Populate the projected matrix */
pcolptr[0] = 0;
for (pnnz=0, ii=0; ii<pncols; ii++) {
    i = cand[ii].val;
    for (j=colptr[ii]; j<colptr[ii+1]; j++) {
        if (rmarker[colind[j]])
            pcolind[pnnz++] = colind[j];
    }

    pcolids[ii] = colids[i];
    pcolptr[ii+1] = pnnz;
}

/* Reset the rmarker array */
if (cid == -1) { /* Initial projection */
    gk_iset(nrows, 0, rmarker);
}
else { /* The other projections */
    for (i=colptr[cid]; i<colptr[cid+1]; i++)
        rmarker[colind[i]] = 0;
}

return pmat;

```

This optimization is essential for sparse data.